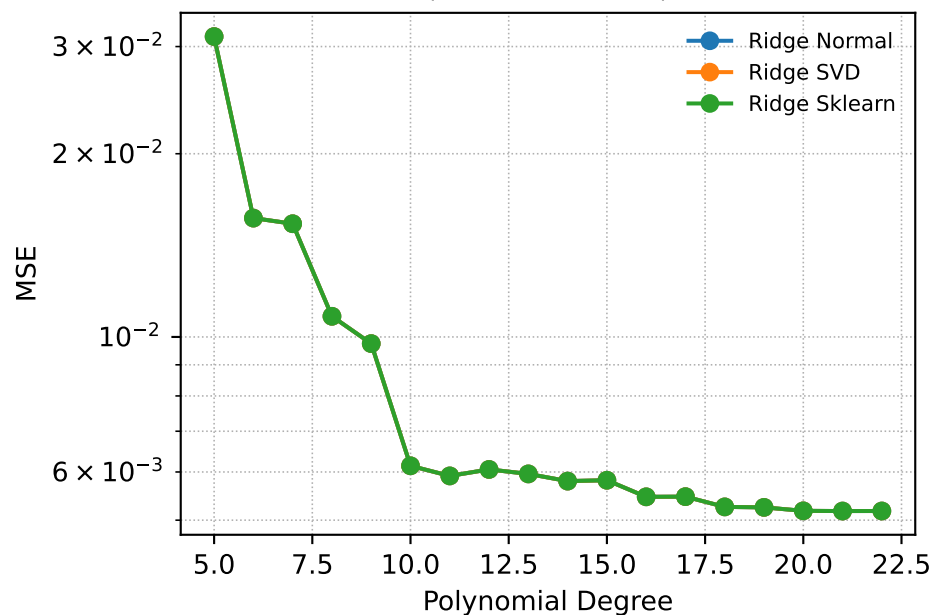
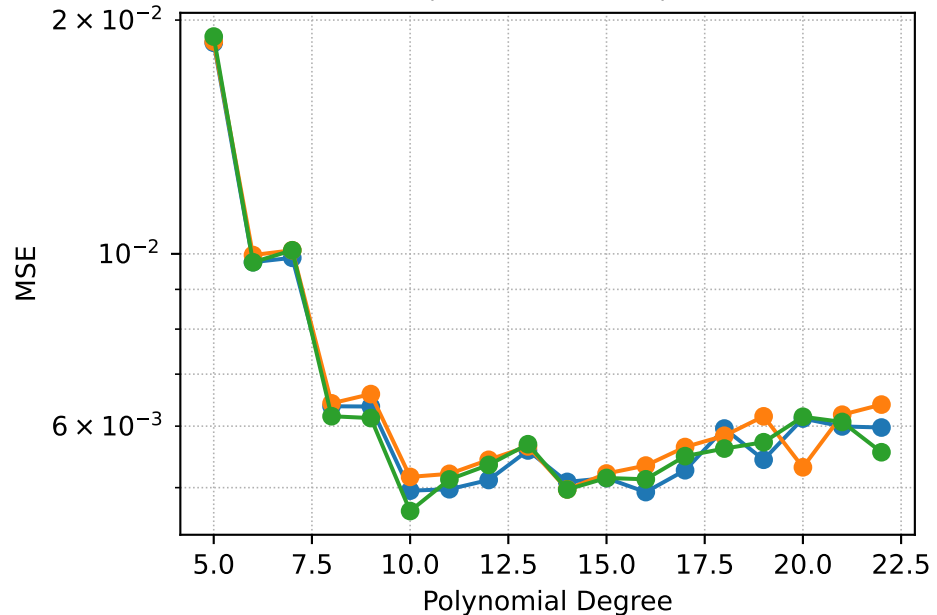


Ridge Regression — Optimal λ per Validation Metric ($n=30, \sigma=0.1$)


Training Loss
(best $\lambda \approx 1.0\text{e-}06$)



Deterministic Holdout
(best $\lambda \approx 1.0\text{e-}06$)



K-Fold CV
(best $\lambda \approx 1.0\text{e-}03$)



Residual Bootstrap
(best $\lambda \approx 1.0\text{e-}06$)

